

# 1 Lamport's Distributed Mutual Exclusion Protocol is Safe

Assume that Lamport's distributed mutual exclusion protocol is not safe. Thus there is an execution of the protocol in which two distinct processes, say  $P_i$  and  $P_j$ , are in their critical sections simultaneously, say at time  $t$  (see Figure 1). Let their critical section requests be denoted by  $R_i$  and  $R_j$ , respectively. Further, let their overlapping critical sections be denoted by  $CS_i$  and  $CS_j$ , respectively. For the critical section  $CS_i$ , we use  $CS_i.begin$  and  $CS_i.end$  to denote its start and finish events. We can likewise define  $CS_j.begin$  and  $CS_j.end$ . Finally, let  $\langle ts_i, i \rangle$  and  $\langle ts_j, j \rangle$  denote timestamps of requests  $R_i$  and  $R_j$ , respectively.

Without loss of generality, assume that  $\langle ts_i, i \rangle < \langle ts_j, j \rangle$ . For  $P_j$  to enter critical section, both L1 and L2 must be true at  $CS_j.begin$ . For L1 to be true at  $CS_j.begin$ ,  $P_j$  must have received a message  $m_i$  from  $P_i$  whose timestamp is greater than  $\langle ts_j, j \rangle$ . There are three different cases depending on when  $P_i$  sends  $m_i$ .

1.  $P_i$  sends  $m_i$  before it generates  $R_i$ : In this case, clearly,  $P_i$ 's logical clock value at the time it sends  $m_i$  is at least  $ts_j$ . Thus,  $P_i$ 's logical clock value at the time it generates  $R_i$  is greater than  $ts_j$ , which contradicts the assumption that  $\langle ts_i, i \rangle < \langle ts_j, j \rangle$ .
2.  $P_i$  sends  $m_i$  after it leaves its critical section: In this case,  $CS_i.end \rightarrow m_i.send \rightarrow m_i.receive \rightarrow CS_i.begin$ . This implies that  $CS_i$  ends before  $CS_j$  starts and the two critical sections do not overlap, which contradicts the assumption that  $P_i$  and  $P_j$  are in their critical sections simultaneously.
3.  $P_i$  sends  $m_i$  after it generates  $R_i$  but before leaves its critical section  $CS_i$ : In this case, at the time  $P_i$  generates  $R_i$ , it will send a REQUEST message to  $P_j$  containing the request  $R_i$ . Since all channels are FIFO,  $P_j$  will receive  $P_i$ 's REQUEST message before it will receive  $m_i$ . Immediately on receiving the REQUEST message,  $P_j$  will insert  $\langle ts_i, i \rangle$  into its priority queue.  $P_i$ 's request will stay in  $P_j$ 's queue at least until time  $t$ . This implies at the time  $P_j$  enters its critical section,  $P_i$ 's request is still in  $P_j$ 's queue, which contradicts with the assumption that L2 is true for  $P_j$  at  $CS_j.begin$ .

Since all three cases lead to a contradiction, we can conclude that Lamport's distributed mutual exclusion protocol is safe.




Figure 1: Two processes  $P_i$  and  $P_j$  are in their critical sections simultaneously.

## 2 Ricart and Agrawala's Distributed Mutual Exclusion Protocol is Safe

Assume that Ricart and Agrawala's distributed mutual exclusion protocol is not safe. Thus there is an execution of the protocol in which two distinct processes, say  $P_i$  and  $P_j$ , are in their critical sections simultaneously, say at time  $t$  (see Figure 2). We use the same notation as before.

Without loss of generality, assume that  $\langle ts_i, i \rangle < \langle ts_j, j \rangle$ . For  $P_j$  to enter critical section,  $P_j$  must have received the REPLY message from  $P_i$  for its request  $R_j$ . There are three different cases depending on when  $P_i$  sends the REPLY message.

1.  $P_i$  sends the REPLY message before it generates  $R_i$ : In this case, clearly,  $P_i$ 's logical clock value at the time it sends the REPLY message is at least  $ts_j$ . Thus,  $P_i$ 's logical clock value at the time it generates  $R_i$  is greater than  $ts_j$ , which contradicts the assumption that  $\langle ts_i, i \rangle < \langle ts_j, j \rangle$ .
2.  $P_i$  sends the REPLY message after it leaves its critical section: In this case,  $CS_i.end \rightarrow REPLY.send \rightarrow REPLY.receive \rightarrow CS_i.begin$ . This implies that  $CS_i$  ends before  $CS_j$  starts and the two critical sections do not overlap, which contradicts the assumption that  $P_i$  and  $P_j$  are in their critical sections simultaneously.
3.  $P_i$  sends the REPLY message after it generates  $R_i$  but before leaves its critical section  $CS_i$ : In this case,  $P_i$  sends the REPLY message while its own request is still pending (not yet completed). This contradicts the fact that, as per the protocol, if  $P_i$  has a pending critical section request, then it cannot send a REPLY message to a request with larger timestamp.

Since all three cases lead to a contradiction, we can conclude that Ricart and Agrawala's distributed mutual exclusion protocol is safe.




Figure 2: Two processes  $P_i$  and  $P_j$  are in their critical sections simultaneously.